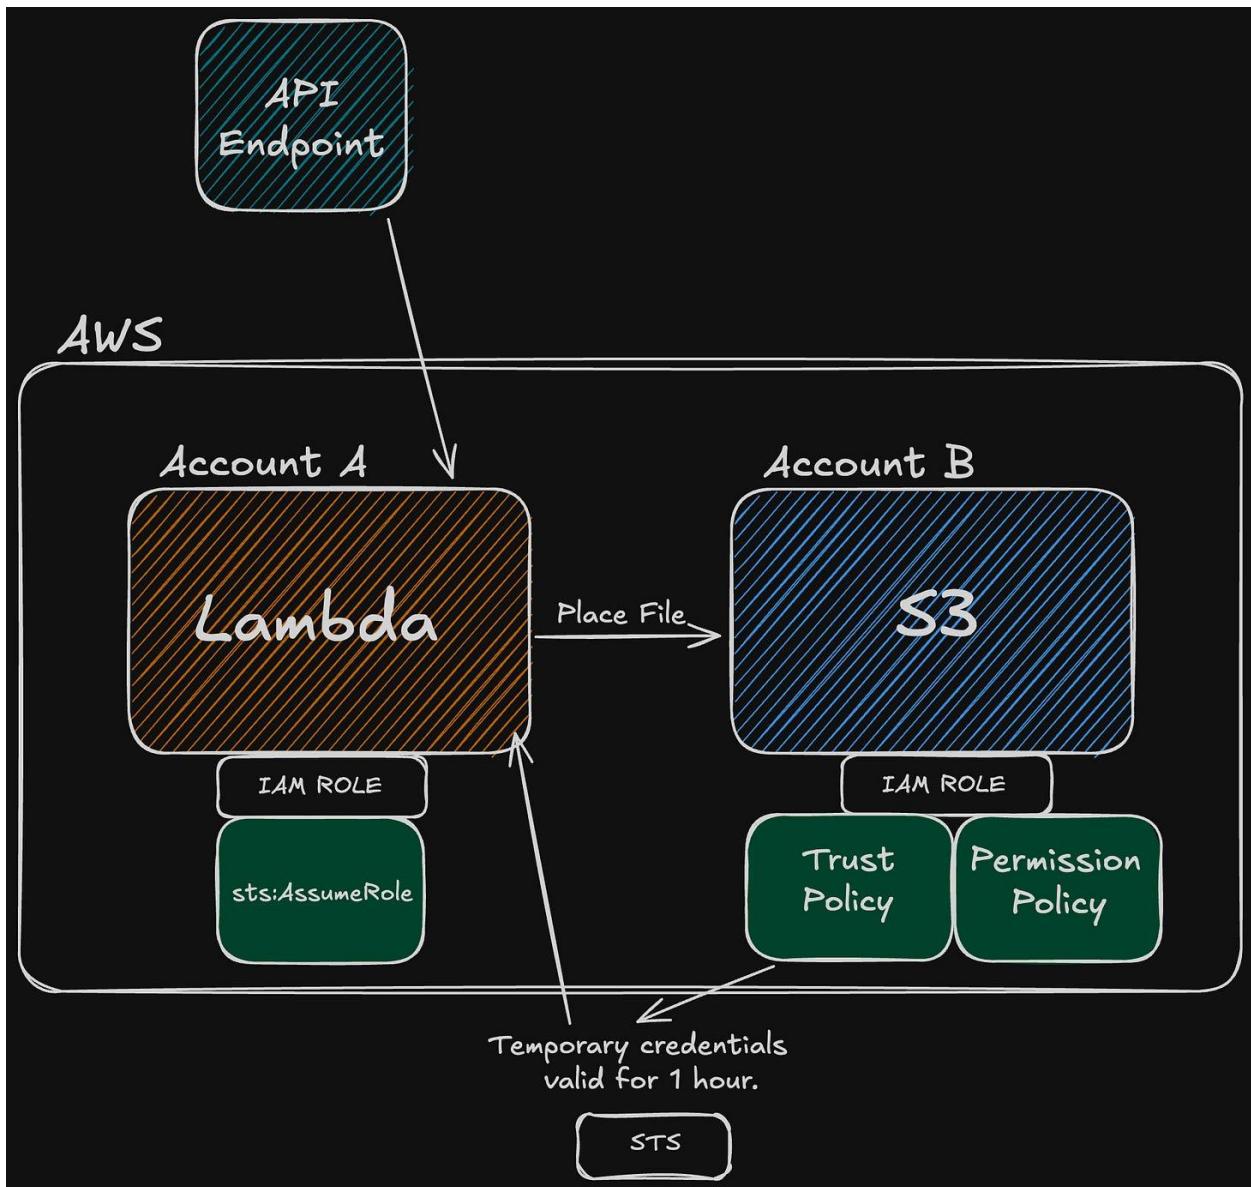


AWS STS

What is it and why is it so important? Okay, basically at a high level, AWS STS lets you request temporary credentials so a service in one account can safely take action another account.

Cross account access gets messy fast if you don't structure it properly, so here's the secure version of how it should work and what we are going to build today.



Account A

Runs something like:

- Lambda
- ECS task

It has a single IAM policy that says: *I am allowed to call sts:AssumeRole on a role in*

Account B

Contains the actual IAM role. This role has two things:

A trust policy that says: *I trust Account A to assume me.*

A permission policy that says: *Here's what I can do in Account B (example: write to S3, deploy to ECS, read Secrets Manager etc).*

Done.

Account B exposes a role.

Account A gets permission to assume it.

This is cool because, it means no permanent credentials and no giant list of access keys floating around.

Project Time

Okay, now you understand what's going on here. Let's build it out in action!

As usual, I reserve the Projects for community members... come join the fun!

If you've made it this far, thank you. Your support genuinely means a lot. And if you're not completely satisfied with this issue, please hit reply and tell me what I can improve. I always want to make this newsletter more useful for you.

Heads up: This project won't fit in your inbox. Click through to view it.

We are going to deploy a Lambda, using Python to call an API Endpoint that tracks the location of the ISS, then dump the JSON results in S3 in another AWS Account.



Step One

Sign into your AWS Account and create a new Lambda. You could deploy this via Terraform. I'm going to use the console since it's best for visual walkthroughs.

Screenshot of the AWS Lambda Functions list:

Functions (2)					
<input type="text"/> Filter by attributes or search by keyword					
<input type="checkbox"/>	Function name	Description	Package type	Runtime	Last modified
<input type="checkbox"/>	cybernotes-LLM	-	Zip	Python 3.13	6 months ago
<input type="checkbox"/>	security-intel-scrapers	-	Zip	Node.js 20.x	4 months ago

A red oval highlights the "Create function" button in the top right corner of the interface.

I'm using Python here, but use whatever you're comfortable with. When prompted, create a new role here:

▼ Change default execution role

Execution role
Choose a role that defines the permissions of your function. To create a custom role, go to the [IAM console](#).

Create a new role with basic Lambda permissions
 Use an existing role
 Create a new role from AWS policy templates

1 Role creation might take a few minutes. Do not delete the role or edit the trust or permissions policies in this role.

Lambda will create an execution role named ISS-Tracker-role-wbjwx6en, with permission to upload logs to Amazon CloudWatch Logs.

Step Two

For this script to work we need the requests library.

We could add the requests library here, so let's quickly do that. In case you didn't know, libraries and external packages need to be separately uploaded to your Lambda as a "layer".

We are actually going to use the built in boto3 URL fetch method. However, I thought it would be cool to quickly show you how to upload a Lambda layer as it's a good skill to have.



The screenshot shows the AWS Lambda Layers interface. At the top, there are tabs for 'Layers' (which is selected) and 'Info'. Below the tabs are two buttons: 'Edit' and 'Add a layer'. There are several filter and sorting options: 'Merge order', 'Name', 'Layer version', 'Compatible runtimes', 'Compatible architectures', and 'Version ARN'. A message at the bottom states 'There is no data to display.'

On your local device, open a terminal and follow these steps:

```
mkdir layers
```

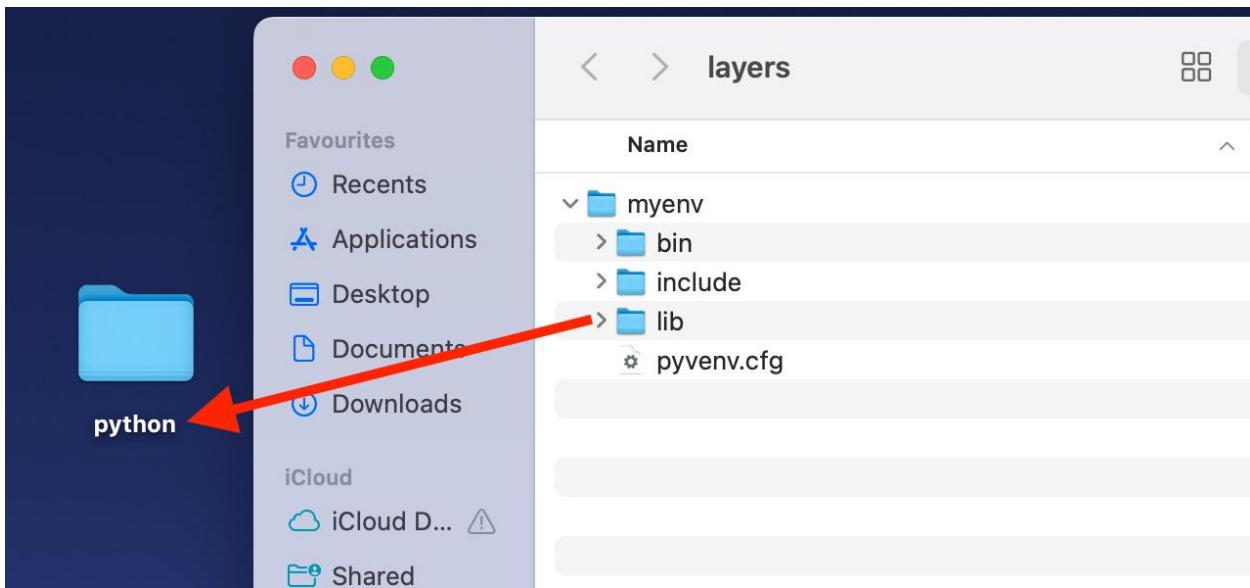
```
cd layers
```

```
python3 -m venv myenv
```

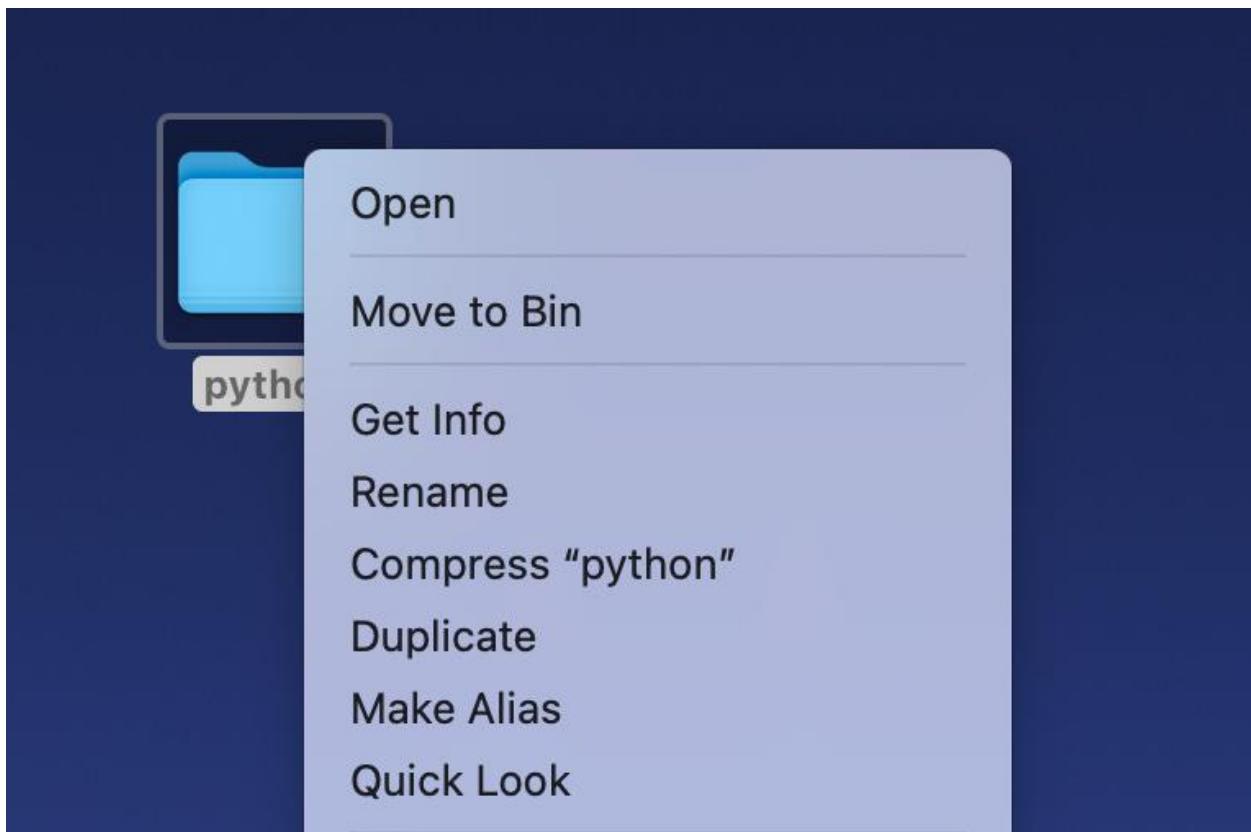
```
source myenv/bin/activate
```

```
pip3 install requests
```

Open this in your finder, create a new folder called python and drag the lib folder over



Compress this file



We can now upload it to our lambda.

Back in the main Lambda page, click into layers and select create layer.

Layer configuration

requests

Description - optional

Cyber Notes Demo

 Upload a .zip file Upload a file from Amazon S3[Choose file](#)

python.zip

3.44 MB

[X](#)

For files larger than 10 MB, consider uploading using Amazon S3.

Compatible architectures - optional | Info

Choose the compatible instruction set architectures for your layer.

Compatible runtimes - optional | Info

Choose up to 15 runtimes.

Python 3.10 [X](#)

Give it a suitable name and upload your compressed Python file, copy the ARN here, you'll need it next.

Add layer**Function runtime settings**Runtime
Python 3.14Architecture
x86_64**Choose a layer****Layer source | Info**Choose from layers with a compatible runtime and instruction set architecture or specify the Amazon Resource Name (ARN) of a layer version. You can also [create a new layer](#). AWS layers

Choose a layer from a list of layers provided by AWS.

 Custom layers

Choose a layer from a list of layers created by your AWS account.

 Specify an ARN

Specify a layer by providing the ARN.

Specify an ARN

Specify a layer by providing the Amazon Resource Name (ARN).

arn:aws:lambda:eu-west-2:202533501269:layer:requests:1

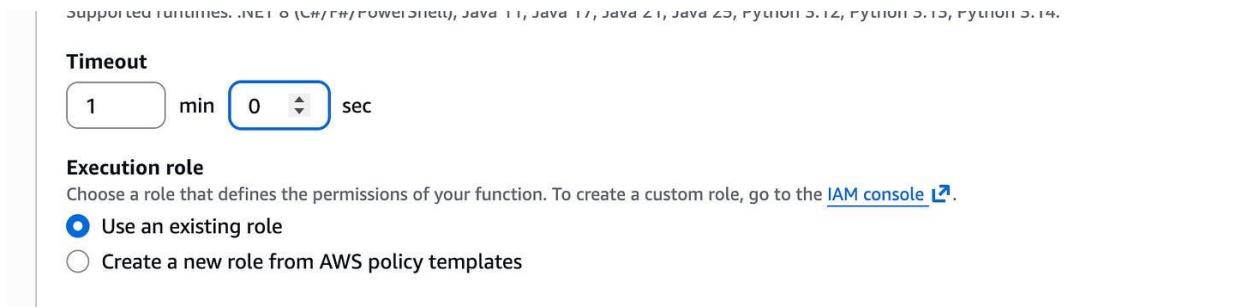
[Verify](#)[Cancel](#)[Add](#)

Go back to your Lambda, add layer and paste your ARN into the box.

Your script would now work using...

```
import requests
```

Layers side track out of the way, change the timeout of your Lambda to 1 min and paste this very small script in.



```
import urllib.request  
  
import json  
  
  
def lambda_handler(event, context):  
    with urllib.request.urlopen("http://api.open-notify.org/iss-now.json") as r:  
        data = json.load(r)  
    return data
```

Deploy and run:

PROBLEMS OUTPUT CODE REFERENCE LOG TERMINAL

Status: Succeeded
Test Event Name: test

Response:

```
{  
  "iss_position": {  
    "latitude": "45.7411",  
    "longitude": "13.7414"  
  },  
  "message": "success",  
  "timestamp": 1763244443}
```

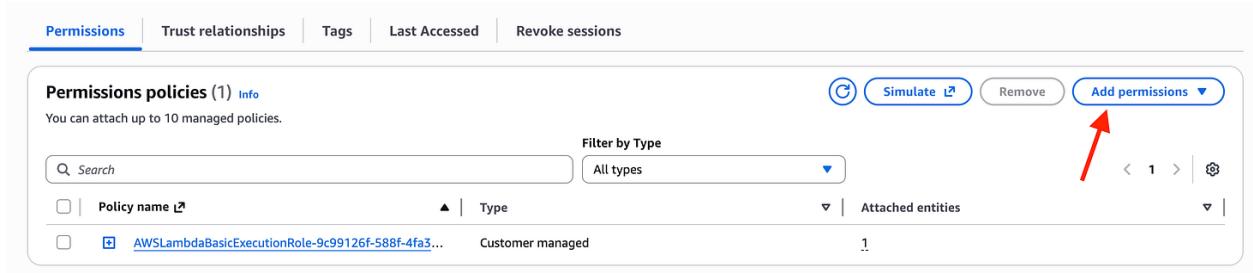
Q

Nice! That's working but we want this Lambda to take the results here and dump them into another AWS account S3 bucket. To do that it needs the permissions we spoke about earlier.

Step Three

Time to configure the roles and policies.

Open the role of your Lambda and attach a policy that gives it sts:AssumeRole into your second AWS Account. Like this:



Just quickly, if you don't have a second AWS account you should make one here. I'm not going to go through that here but it's pretty straight forward and can be done under AWS Organisations, you'll just need a second email address also.

Create an inline policy

The screenshot shows the AWS IAM Policy editor interface. The left panel displays the JSON code for an inline policy:

```
1▼ {  
2    "Version": "2012-10-17",  
3    "Statement": [  
4        {  
5            "Effect": "Allow",  
6            "Action": "sts:AssumeRole",  
7            "Resource": "arn:aws:iam::<PASTEYOURACCOUNTIDHERE>;role/iss_lambda"  
8        }  
9    ]  
10 }
```

The line number 7 is highlighted with a red box. The right panel contains sections for managing statements, adding actions, choosing services, and adding resources. At the bottom, there are status indicators for security, errors, warnings, and suggestions.

You'll want to give it the permissions here to assume the role we're going to deploy into your second account with the name and ID as the resources.

Log out of this account and log into your second new account.

Create a new IAM role and select “Custom Trust Policy”.

Select trusted entity Info

Trusted entity type

AWS service

Allow AWS services like EC2, Lambda, or others to perform actions in this account.

AWS account

Allow entities in other AWS accounts belonging to you or a 3rd party to perform actions in this account.

Web identity

Allows users federated by the specified external web identity provider to assume this role to perform actions in this account.

SAML 2.0 federation

Allow users federated with SAML 2.0 from a corporate directory to perform actions in this account.

Custom trust policy

Create a custom trust policy to enable others to perform actions in this account.

Paste this in to allow Account A to assume the role here.

```
{
```

```
    "Version": "2012-10-17",
```

```
    "Statement": [
```

```
        {
```

```
            "Effect": "Allow",
```

```
            "Principal": {
```

```
                "AWS": "arn:aws:iam::<ACCOUNT_A_ID>:root"
```

```
            },
```

```
            "Action": "sts:AssumeRole"
```

```
        }
```

```
    ]
```

```
}
```

Attach the permissions your Lambda will need to **PUT** a file inside the S3 bucket.

```
{
```

```
    "Version": "2012-10-17",
```

```
    "Statement": [
```

```
        {
```

```

    "Effect": "Allow",
    "Action": [
        "s3:PutObject",
        "s3:PutObjectAcl"
    ],
    "Resource": "arn:aws:s3:::YOUR_BUCKET_NAME/*"
}
]
}

```

Paste this in to allow Account A to assume the role here.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "AWS": "arn:aws:iam::<ACCOUNT_A_ID>:root"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

Attach the permissions your Lambda will need to **PUT** a file inside the S3 bucket.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "s3:PutObject",
                "s3:PutObjectAcl"
            ],
            "Resource": "arn:aws:s3:::YOUR_BUCKET_NAME/*"
        }
    ]
}
```

We can update this in a second when we've made our bucket.

Navigate to S3, create a bucket with the basic configuration, note the name and update your role:

The screenshot shows the AWS S3 console interface. At the top, there are two tabs: 'General purpose buckets' (which is selected) and 'All AWS Regions'. Below the tabs, there's a search bar labeled 'Find buckets by name' and a filter section with dropdowns for 'Name' (set to 'isslocation'), 'AWS Region' (set to 'Europe (Stockholm) eu-north-1'), and 'Creation date' (set to 'November 15, 2025, 22:49:26 (UTC+00:00)'). On the right side of the search bar are navigation arrows and a settings gear icon. Above the search bar, there are four buttons: a blue circular button with a 'C' (Copy ARN), a grey 'Empty' button, a grey 'Delete' button, and a yellow 'Create bucket' button. The main list area shows one item: 'General purpose buckets (1) [Info](#)'. Below this, a note says 'Buckets are containers for data stored in S3.' The single bucket entry is 'isslocation'.

Step Four

Finally, let's go back to our Account A and update our script to assume the role in Account B and dump the JSON into the bucket we just made.

```
import urllib.request  
import json  
import boto3  
from datetime import datetime
```

```
# Config  
ROLE_ARN = "arn:aws:iam::<ACCOUNTB>:role/iss_lambda"  
BUCKET_NAME = "isslocation"
```

```
def lambda_handler(event, context):  
    # 1. Fetch ISS JSON  
    with urllib.request.urlopen("http://api.open-notify.org/iss-now.json") as r:  
        data = json.load(r)
```

```
# 2. Assume role in Account B  
  
sts_client = boto3.client("sts")  
  
assumed_role = sts_client.assume_role(  
    RoleArn=ROLE_ARN,  
    RoleSessionName="LambdaToS3Session"  
)
```

```
credentials = assumed_role['Credentials']
```

```
# 3. Create S3 client using assumed role credentials  
  
s3_client = boto3.client(  
    "s3",  
    aws_access_key_id=credentials['AccessKeyId'],  
    aws_secret_access_key=credentials['SecretAccessKey'],  
    aws_session_token=credentials['SessionToken'])
```

```
# 4. Generate a filename based on timestamp
```

```
timestamp = datetime.utcnow().strftime("%Y-%m-%dT%H-%M-%SZ")  
  
s3_key = f"iss-location-{timestamp}.json"
```

```
# 5. Upload to S3
```

```
s3_client.put_object(  
    Bucket=BUCKET_NAME,  
    Key=s3_key,
```

```
Body=json.dumps(data),  
ContentType="application/json"  
)  
  
return {"status": "success", "s3_key": s3_key}
```

Finally, let's go back to our Account A and update our script to assume the role in Account B and dump the JSON into the bucket we just made.

```
import urllib.request
import json
import boto3
from datetime import datetime

# Config
ROLE_ARN = "arn:aws:iam::<ACCOUNTB>:role/iss_lambda"
BUCKET_NAME = "isslocation"

def lambda_handler(event, context):
    # 1. Fetch ISS JSON
    with urllib.request.urlopen("http://api.open-notify.org/iss-now.json") as r:
        data = json.load(r)

    # 2. Assume role in Account B
    sts_client = boto3.client("sts")
    assumed_role = sts_client.assume_role(
        RoleArn=ROLE_ARN,
        RoleSessionName="LambdaToS3Session"
    )

    credentials = assumed_role['Credentials']

    # 3. Create S3 client using assumed role credentials
    s3_client = boto3.client(
        "s3",
        aws_access_key_id=credentials['AccessKeyId'],
        aws_secret_access_key=credentials['SecretAccessKey'],
        aws_session_token=credentials['SessionToken']
    )

    # 4. Generate a filename based on timestamp
    timestamp = datetime.utcnow().strftime("%Y-%m-%dT%H-%M-%SZ")
    s3_key = f"iss-location-{timestamp}.json"

    # 5. Upload to S3
    s3_client.put_object(
        Bucket=BUCKET_NAME,
        Key=s3_key,
        Body=json.dumps(data),
        ContentType="application/json"
    )

    return {"status": "success", "s3_key": s3_key}
```

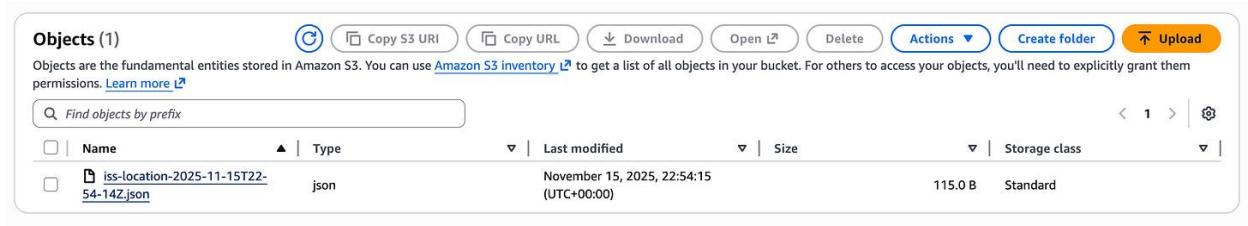
Test the script

```
PROBLEMS    OUTPUT    CODE REFERENCE LOG    TERMINAL

Status: Succeeded
Test Event Name: test

Response:
{
  "status": "success",
  "s3_key": "iss-location-2025-11-15T22-54-14Z.json"
}
```

Nice! Let's go check the bucket:



The screenshot shows the AWS S3 'Objects (1)' page. At the top, there are several buttons: Copy S3 URI, Copy URL, Download, Open, Delete, Actions (with a dropdown arrow), Create folder, and Upload. Below these are buttons for Find objects by prefix and a search bar. A navigation bar shows page 1 of 1. The main table lists one object:

Name	Type	Last modified	Size	Storage class
iss-location-2025-11-15T22-54-14Z.json	json	November 15, 2025, 22:54:15 (UTC+00:00)	115.0 B	Standard

Congrats, you've just learnt one of the most important Cloud Security concepts! 🎉

Remember this is really helpful because you don't necessarily need to own both the accounts you're setting this up with. What if you're using a product that's hosted in an external AWS account not owned by you? You should never give the vendor access keys and you decide exactly what actions they get.